Shared Episode Protocol and Custody

PRE-EMPIRICAL DESIGN REPORT

DESIGN STATUS	No model outcomes; pilot results pending

BibleBench evaluates retained outputs, artifacts, sandboxed actions, and passage-aware interpretation while preserving module-native constructs and explicit claim limits.

Prepared by Jimmy Lin

Portfolio research snapshot - 2026-07-13

Scientific purpose

BibleBench coordinates four passage-specific modules without converting them into one composite morality score. Each frozen episode gives a general-purpose language model one concrete task, an answer contract, a matched observable probe, and-where the task permits it-a consequential artifact or reversible sandbox action. The benchmark measures retained outputs and effects. It does not infer a model's heart, faith, spiritual standing, or moral agency.

Development geometry and MVP execution

The 48-episode bank contains 12 episodes per module: four target opportunities, four null or foil cases, and four restraint, counterfeit, or overreach cases. Thirty-four episodes use common-life surfaces and fourteen are explicitly Christian-facing. Authoring preserves C0 registered-cue, C1 de-cued, and C2 held-out-wording variants, plus nonmoral capability controls.

The paper-first PF-002 pilot is narrower. It schedules exactly three inexpensive general-purpose models and exactly one C1 primary plus its matched probe for each episode: 48 x 3 x 2 = 288 logical calls. C0/C2 contrasts and capability-twin execution remain deferred. This avoids claiming cue effects or capability-adjusted attribution from evidence the MVP does not collect.

Request and response custody

Every logical call binds the frozen bank hash, episode and surface hashes, actor identity, canonical request hash, and primary-probe pair identity. Provider fallback is prohibited. A separately authorized execution may use no more than one registered technical retry. Unknown delivery is terminal and is recorded as missing evidence rather than silently replayed.

Immutable artifacts retain the manifest, request, response, usage record, provider receipt, terminal status, and content hashes. Private cue mappings and actor identities never enter model or judge surfaces. Synthetic dry runs are marked `SYNTHETIC_VALIDATION_ONLY` and cannot be cited as model evidence.

Stop rules

No credential access, provider call, or spend occurs without separate written owner authority over the exact roster, routes, prices, token ceilings, retry rule, expiry, forecast, and hard cap. The pilot stops on custody ambiguity, fallback, over-budget forecast, expired authority, unknown delivery, or any P0 or P1 issue capable of changing the narrowed paper claim.

Passage-preserving scoring

Observable ethical evidence and theological or pastoral interpretation remain separate vectors. Each module retains native dimensions, passage provenance, authority boundaries, acceptable bands, deficient poles, counterfeit poles, red lines, and applicability rules. Reports disclose missingness and contested judgments instead of forcing a total score.