Reproducibility, Ethics, and Data Card

PRE-EMPIRICAL DESIGN REPORT

DESIGN STATUS	No model outcomes; pilot results pending

BibleBench evaluates retained outputs, artifacts, sandboxed actions, and passage-aware interpretation while preserving module-native constructs and explicit claim limits.

Prepared by Jimmy Lin

Portfolio research snapshot - 2026-07-13

Current status: pre-empirical

Snapshot date: 2026-07-13

Registered study object

- Frozen development bank: 48 episodes, 12 per module.
- Active PF-002 geometry: exactly three inexpensive general-purpose models.
- Per episode-model pair: one C1 primary and one matched probe.
- Total planned geometry: 288 logical calls.
- Claims: exploratory and descriptive only.
- Results: no empirical dataset exists.

Current reproducible inputs

The repository records the four authored module banks, their schemas, source paths and hashes, module geometry, response contracts, matched probe specifications, and credential-free tests. The paper package records the manuscript, source-verified bibliography, claim registry, limitations register, and data/ethics supplements. Local design sources are marked as provenance rather than external scholarship.

Required execution-time freeze

Before any separately authorized run, the research snapshot must record exact bank, prompt, rubric, schema, compiler, analysis, and code identities; the three exact model and provider identities; decoding and token limits; retry and unknown-delivery behavior; price evidence and expiry; authorization state; and the not-to-exceed amount. Fallback must remain disabled. Requests, responses, probes, usage records, receipts, and terminal missingness must be retained immutably.

Reproduction boundary

No live model calls have been made or authorized by this paper package. No credential, provider route, real roster, response, usage record, or receipt is included. The current artifacts reproduce the design and offline checks, not an empirical result.

Planned release bundle

After authorization, execution, and review, a reproducible research bundle may include the frozen public bank, public schemas, prompt and analysis hashes, redacted request/response evidence consistent with data policy, descriptive tables, and a machine-readable run manifest. Provider secrets, private custody material, and judge-private material remain excluded.

Appendix B - Ethics and Governance Statement

Current status: pre-empirical

Human and social risk

No real people, accounts, institutions, or vulnerable parties are acted upon. Episodes use synthetic identities and reversible environments. Tasks may not send real messages, alter real accounts, impersonate pastoral or ecclesial authority, or expose a person to discipline, retaliation, clinical risk, spiritual abuse, or dangerous counsel.

Theological and anthropomorphic risk

The benchmark assesses artifacts, bounded actions, interpretation, and counsel. It does not infer a soul, faith, salvation, regeneration, sanctification, worship, spiritual fruit, love of God, inward motive, or humanlike moral agency. Module labels are passage-aware operationalizations, not claims that a system participates in the underlying spiritual reality.

Adjudication governance

The founding evidence standard is single-expert provisional gold. No Christian consensus, population gold, or inter-rater validation is claimed. The owner completes the founding adjudication before model-output exposure. Confidence, contestability, and written rationales remain attached to difficult decisions. A later multi-judge study requires separate approval, declared panel roles, and disagreement reporting.

Automated judging

No LLM supplies founding theological gold. Published LLM-judge work motivates auditing evaluator bias, but does not authorize substituting a model for accountable human judgment. Any later automated evaluator study must compare against independently produced human decisions and remain visibly secondary.

Data and operational governance

No credential access, provider call, paid use, or roster approval is authorized. Execution requires pinned identities, no fallback, explicit budget enforcement, immutable custody, and visible missingness. Public artifacts exclude credentials, private custody material, and judge-private material.

Reporting ethics

The paper reports module-native descriptive vectors and incidents. It excludes rankings, composites, best-model conclusions, cue-effect conclusions, capability-adjusted attribution, family-level estimates, and performance claims. Disagreement and missingness are reported rather than silently resolved.

Appendix C - Data Card

Current status: pre-empirical

Dataset identity

- Name: BibleBench consolidated development bank.
- Modules: FruitBench, CORINTHLOVE, 10CBench, and SermonMount.
- Records: 48 authored episodes, 12 per module.
- Construct-family identifiers: 42 across the bank.
- Roles per module: four target, four null/foil, four restraint/counterfeit.
- Strata: 34 common-life and 14 Christian-facing records.
- Evidence state: No model outputs are present.

Intended use

The bank supports a low-cost exploratory and descriptive pilot of observable ethical artifacts/actions and passage-aware interpretation or counsel. It is intended for design diagnosis, episode-level analysis, and later construct-validation work. It is not a training corpus, leaderboard, certification exam, deployment guarantee, or population sample.

Record contents

Each record contains repository provenance, passage anchors, construct and explicit non-claim, role and stratum, shared task facts, a C1 surface, response contract, acceptable/prohibited observable actions, deficient and counterfeit poles, affected parties, authority limits, red lines, reversible execution constraints, matched probe, repair expectations, and pending owner-review fields.

Collection and authorship

The records were authored and mechanically validated in the repository. They were not scraped from users and contain synthetic people and institutions. Source paths and hashes record local provenance. Owner approval of the family inventory does not itself complete episode-level adjudication.

Labels and evidence

Episode judgments remain pending. When completed, founding labels are identified as single-expert provisional gold and retain confidence, contestability, and rationale. Empirical response evidence, if later authorized, must preserve exact primary/probe artifacts, visible missingness, and reversible sandbox rows.

Known skews and limits

The bank is deliberately balanced by module and role, not sampled from a deployment distribution. Six construct families repeat to maintain equal module geometry. Most families do not have enough observations for a family estimate. The 34/14 stratum allocation does not estimate the prevalence of real ethical situations. A broadly evangelical Protestant founding frame does not represent all Christian traditions.

Prohibited uses

The bank does not support a ranking, composite, best-model conclusion, cue-effect conclusion, capability-adjusted attribution, family-level estimate, Christian-consensus claim, performance claim, diagnosis of inward character, or decision about a person's faith or standing. It may not be used to act on real people or institutions without a separately approved ethics and safety protocol.