Analysis Tables, Figures, and Claims Boundary

PRE-EMPIRICAL DESIGN REPORT

DESIGN STATUS	No model outcomes; pilot results pending

BibleBench evaluates retained outputs, artifacts, sandboxed actions, and passage-aware interpretation while preserving module-native constructs and explicit claim limits.

Prepared by Jimmy Lin

Portfolio research snapshot - 2026-07-13

Registered reporting stance

The first pilot is exploratory and descriptive. Every pre-run table and figure contains `PENDING_EMPIRICAL_RUN`; synthetic fixtures are plumbing tests only. The report presents vectors and missingness rather than a composite, best-model claim, or leaderboard.

Table shells

1. Completion, refusal, malformed output, truncation, rate limits, unknown

delivery, and terminal missingness by actor and module.

2. Module-native dimension vectors with count, denominator, missing count, and

founding-judge confidence.

3. Common-life and Christian-facing stratum summaries.
4. Primary-action and matched-probe consistency.
5. Red-line incidents with episode, evidence availability, and contestability.
6. Single-expert decision, confidence, contestability, and blinded-repeat

stability summaries.

Figure contracts

Figures receive tidy rows from the registered JSON contracts. Permitted marks include completion bars, module-native dot plots, stratum facets, action-probe consistency matrices, red-line incident counts, and confidence or contestability distributions. Missing values remain visible and do not become zeros.

Explicit exclusions

There is no leaderboard, qualified composite, best-model field, p-value, family-level estimate, cue effect, or capability-adjusted attribution in the PF-002 report. C0/C2 and capability twins are not executed. One judge cannot support Christian-consensus or inter-rater claims. Forty-eight development episodes cover interfaces and construct families but do not estimate every family independently.

Promotion rule

Only actual, custody-valid model observations and frozen owner judgments may replace the empirical sentinel. A failed or incomplete run is reported as a failure account. It is never repaired with synthetic values or undocumented replays.

Registered figure data

- completion_missingness: bar - fields: actor, module, status, count, denominator
- module_native_vectors: dot - fields: actor, module, dimension, value, missing_count
- strata: facet_bar - fields: actor, module, stratum, count, denominator
- action_probe_consistency: matrix - fields: actor, module, consistency, count
- red_line_incidents: bar - fields: actor, module, incident_class, count, evidence_available
- judge_confidence_contestability: stacked_bar - fields: confidence, contestability, count

Claims registry

D001 [DESIGN_CLAIM]: BibleBench assesses observable outputs, consequential artifacts, reversible sandbox actions, registered omissions, downstream effects, and passage-aware interpretation or counsel; it does not establish inward spiritual states or Christian moral agency.

Boundary: Measurement scope only; no model has been evaluated.

D002 [DESIGN_CLAIM]: The shared design preserves FruitBench, CORINTHLOVE, 10CBench, and SermonMount as passage-specific modules rather than merging them into one undifferentiated Christian-ethics construct.

Boundary: Shared mechanics do not make the module-native constructs interchangeable.

D003 [DESIGN_CLAIM]: The frozen development bank contains 48 episodes, 12 per module, with four target, four null or foil, and four restraint or counterfeit roles per module, and 34 common-life plus 14 Christian-facing records overall.

Boundary: Coverage is not a family-level estimate or a population sample.

D004 [DESIGN_CLAIM]: Under PF-002, the active low-cost MVP contains exactly three inexpensive general-purpose models and 48 episodes, yielding 288 logical calls when each model receives one C1 primary and one matched probe per episode.

Boundary: Geometry is approved; exact roster, provider routes, calls, and spend are not.

D005 [DESIGN_CLAIM]: The active MVP gives each episode-model pair exactly one C1 primary and one immediately matched observable probe.

Boundary: Earlier C0/C1/C2 and forecast geometry is deferred design history, not active execution wording.

D006 [DESIGN_CLAIM]: The MVP report is exploratory, descriptive, and vector-first and excludes rankings, composites, best-model, cue-effect, capability-adjusted, family-level, Christian-consensus, and performance claims.

Boundary: No model outcome or comparative conclusion exists.

D007 [DESIGN_CLAIM]: The founding adjudication protocol uses one theologian-technical expert and stores accepted judgments as single-expert provisional gold, with a blinded repeat frozen before pilot-output exposure.

Boundary: No episode gold is complete and no Christian consensus is claimed.

P001 [PROVISIONAL_THEOLOGICAL_JUDGMENT]: The founding frame treats Scripture as primary, uses a broadly evangelical Protestant public core, uses the Westminster Standards as a calibration lens rather than an identity gate, and retains contested applications.

Boundary: Owner-approved founding judgment; not an ecumenical or population consensus.

P002 [PROVISIONAL_THEOLOGICAL_JUDGMENT]: Godward, covenantal, worship-related, and Spirit-wrought teachings are not behaviorally attributable to a language model merely from its outward output.

Boundary: Such content may be assessed only as interpretation, counsel, or role-boundary fidelity.

X001 [EMPIRICAL_CLAIM_PENDING]: An authorized pilot would describe delivery, missingness, module-native vectors, common-life and Christian-facing strata, action/probe consistency, red-line incidents, and adjudication uncertainty.

Boundary: No ranking, composite, cue effect, performance, or causal claim may be promoted.

X002 [EMPIRICAL_CLAIM_PENDING]: The founding repeat may estimate intra-rater stability only; a later preregistered multi-judge study is required for inter-rater or cross-panel disagreement claims.

Boundary: A single founding expert cannot establish Christian consensus.