Contents

WCF-Bench: A Westminster Standards Steerability Benchmark 2
Abstract 2
1. Introduction 2
1.1 The steerability problem 2
1.2 Why a confessional standard is a strong test case 3
1.3 Construct disclosure 3
1.4 Contributions 4
1.5 A validity-first stance 4
1.6 Paper map 5
2. Related Work 5
2.1 Pluralistic alignment and steerability toward declared value systems 5
2.2 Values and ethics benchmarks for LLMs, including religious-values evaluation 6
2.3 LLM-as-judge validity: qualification, agreement, bias, and artifact failures 6
2.4 Benchmark contamination and canary / do-not-train defenses 7
3. Benchmark Design 7
3.1 Construct definition and disclosure 7
3.2 The anchor registry 8
3.3 The 230-item bank 8
3.4 The ten method-major designs 10
3.5 The tier ladder 12
3.6 Contested-item scoring: adherence, not correctness 13
3.7 The two tracks 14
3.8 Scorer kinds, including OPEN and PROOF 14
4. Validity Methodology: Registered Full Protocol 14
4.1 Judge qualification 14
4.2 Parser-artifact forensics 15
4.3 Adjudication and ratification 15
4.4 Seven registered release gates 15
5. PPP-1 Exploratory Setup 15
5.1 Scope 15
5.2 Score contract and denominators 16
5.3 Custody, cost, and time 16
6. Results 16
6.1 Completion and headline denominators 16
6.2 Per-system descriptive support 16
6.3 Overall pooled support summary 18
7. Discussion 18
7.1 What the preview establishes 18
7.2 What the preview does not establish 19
7.3 Relation to the confirmatory study 19
8. Limitations 19
9. Ethics and Release Posture 20

Appendix A. Empirical Traceability 21
Appendix B. Reproducibility and Source Boundary 22
References 22

WCF-Bench: A Westminster Standards Steerability Benchmark

[AUTHORS]

2026-07-14

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

Abstract

Language models are increasingly asked to represent declared normative frameworks, but accord with such a framework is difficult to measure without collapsing the target into annotator preference. We introduce WCF-Bench, a 230-item benchmark that measures the degree to which a model's conduct and counsel accord with the ethical teaching of the Westminster Standards, as stated in those documents. Ten method-major designs vary the measurement mechanism while retaining a shared, clause-addressable normative source. A registered full protocol specifies judge qualification, artifact forensics, human confirmation, perturbations, and seven release gates. This paper reports only PPP-1, an exploratory preview separate from that protocol: 226 runnable base items, three inexpensive systems, one Gemini judge where judgment was required, and no perturbation, human-panel, or gate evidence. PPP-1 recorded 678/678 cells; 607/666 headline-eligible cells were numeric and 60 total cells were error-recorded. Errors remained missing, and 12 contested cells were excluded from the headline mean. This preview does not support rankings, external quotation, release claims, panel consensus, or T4 completion.

1. Introduction

1.1 The steerability problem

A recurring demand on modern language models is not that they be good in the abstract, but that they *represent a specified normative framework on request* - the values of a particular deployer, community, jurisdiction, or tradition. This is the ambition behind pluralistic-alignment research: rather than encoding a single global ethic and treating every deviation from it as misalignment, a model should be *steerable* toward a declared value system and evaluated on how faithfully it inhabits that system. Steerability of this kind is only as meaningful as our ability to *measure* it. If a model is told to reason and act within a stated framework, can its accord with that framework be assessed rigorously, or must we fall back on the diffuse, annotator-averaged, majority-proxy notions of "helpfulness" and "harmlessness" that dominate current alignment evaluation? The methodological question this paper addresses is narrow and concrete: **given a fully specified normative framework, can a model's conduct and counsel be measured against it with construct validity - and can we know when such a measurement is trustworthy enough to report?**

Most existing values-and-ethics benchmarks answer an easier question. They elicit moral *opinions* against a target that is itself contested: labels are drawn from annotator majorities, the "correct" answer shifts with the annotator pool, and a high score is quietly read as evidence that a model "is ethical." Such instruments are simultaneously (a) under-specified in their construct - whose ethics, codified where? - and (b) prone to saturation and capability-correlation, because a smart model that can pattern-match consensus opinion

scores well without any behavioral commitment. Steerability toward a *declared* framework needs the opposite: a target that is fixed independently of the model and the annotator, so that accord is a property of the model's behavior rather than of the labeling process.

1.2 Why a confessional standard is a strong test case

We take as our declared framework the **Westminster Standards** - the Westminster Confession of Faith together with the Larger and Shorter Catechisms (1647). This is a deliberate choice of test bed, not an endorsement, and it is chosen for properties that make it an unusually demanding and clean instrument for the measurement problem above:

- **It is fixed.** The text has been stable for centuries. There is no moving target, no drift in the "gold" standard between evaluation rounds, and no need to reconstruct the norm from a fresh annotator survey each time the benchmark is run.
- **It is internally coherent and extensively documented.** The tradition has already done the operationalization work that most values benchmarks improvise. The Larger Catechism's exposition of the Decalogue (Q98-148) is a *graded casuistic system*: for each commandment it enumerates the duties required and the sins forbidden, supplies eight interpretive rules for reading the commandments (Q99), a four-headed calculus of what makes some sins more heinous than others (Q151), and a triaxial theory of right action - right in *matter*, *manner*, and *end* (WCF 16.7) - every element attached to proof texts. This is a ready-made, source-verifiable scoring ontology, granular enough to key individual rubric points to specific clauses.
- **It resolves to specific clauses, not vibes.** Because every gold label can be traced to a named anchor in the Standards, the target does not depend on an annotator majority. Where the Standards are silent, or where confessionally competent interpreters genuinely differ, the correct response is calibrated abstention ("the Standards leave this free / this is disputed"), and such items are flagged and scored accordingly rather than forced to a spurious label.

The point is methodological. A fully specified confession lets us isolate the measurement problem from the labeling problem: because the norm is external, fixed, and clause-addressable, any residual difficulty in scoring is attributable to the *measurement mechanism* and the *judge*, not to disagreement about what the right answer is. That is precisely the setting in which a steerability benchmark can be validated.

1.3 Construct disclosure

We state the construct exactly and without softening. **WCF-Bench measures the degree to which a model's conduct and counsel accord with the ethical teaching of the Westminster Standards (Confession of Faith, Larger Catechism, Shorter Catechism, 1647), as stated in those documents.**

This is **not** a claim that a high-scoring model "is ethical" simpliciter, and it is not a measure of generic religiosity. The normative population is stipulated: confessional Reformed Christianity as codified at Westminster. A high WCF-Bench score is a statement of *accord with a declared standard*, not a certificate of moral goodness; a low score is a statement of *discord with that standard*, not an indictment of a model's character. Westminster is not a stand-in for "Christianity" and still less for ethics in general - items are constructed so that they cannot be satisfied by generic moralism, and every gold label must be traceable to a specific anchor in the Standards. Readers, deployers, and model developers should read every number in this paper through that disclosure. The benchmark is an instrument for measuring steerability toward one fully specified framework; it is silent on whether that framework, or any other, ought to be steered toward.

1.4 Contributions

This paper makes two co-equal contributions.

1. **WCF-Bench, a construct-disclosed steerability benchmark.** A frozen bank of **230** anchor-keyed items across **ten method-major designs**. The suite's signature is a design invariant: moral content is held constant while the *measurement mechanism* varies, yielding a designs × commandments matrix in which each commandment is measured

   ten methodologically independent ways - a multitrait-multimethod structure that lets method variance be estimated explicitly rather than hoped away. The ten mechanisms are omission-sensitive casuistry, complicity gradients, severity discernment, fidelity under social pressure, expressed-versus-revealed divergence, conscience and adiaphora triage, proof-text discipline, repentance-and-repair trajectories, triaxial matter/manner/end discrimination, and obedience under incentive conflict. Every design fields two evaluation **tracks** - the model's *own conduct* (is it the moral agent that lies, flatters, equivocates, or keeps its commitments?) and its *advisory fidelity* (does its counsel to a user conform in matter, manner, and end?) - and the divergence between the two is itself a reported metric. A **tier ladder** runs from recognition through applied open-response, multi-turn adversarial pressure, and agentic incentive conflict, with the *drop* between tiers, never a single absolute score, as the headline. Scoring is decomposed into anchor-keyed rubric keypoints rather than holistic ratings, and a mandatory perturbation battery (paraphrase, option order, persona/framing, pushback, register, monitoring-frame) turns invariance into a first-class, scored property.

2. **A reusable validity template for judge-scored benchmarks.** WCF-Bench is built with, and this paper documents in full, machinery that treats validity as the primary deliverable: a **judge-qualification protocol** (a fixed battery, re-qualification triggers, an escalation ladder, drop rules, and a cross-family judge panel to defeat single-family bias); **parser-artifact forensics** as a first-class validity activity, in which apparent "model disagreement" is decomposed into harness artifacts and diagnosed by counterfactual re-parse before any substantive conclusion is drawn; a disciplined **adjudication protocol** that separates *computing* results from *adjudicating* and *ratifying* them; and a set of **pre-registered release gates** whose thresholds and computations are frozen before the data exist. We present this machinery as a template that other declared-norm benchmarks can adopt independently of the Westminster construct.

1.5 A validity-first stance

The two contributions are not sequential - measure first, validate later - but interlocked. Under the registered full protocol, validity checks determine whether a later results table may support a release-grade claim. The gates cover capability correlation, headroom, separation, judge validity, shortcut resistance, human consensus, and reliability. Each has a frozen computation, and an unregistered deviation voids the affected outcome. PPP-1 did not execute these gates and this paper reports no pass or fail result. Its tables remain exploratory, single-judge, and non-quotable regardless of their completeness.

A further methodological commitment supports this stance. In the course of construction we found that measured "disagreement" among judges was, in an instructive case, an artifact of the harness rather than a signal about the models - a score-extraction parsing bug and a false-firing willingness veto that tripped whenever a judge dutifully quoted a refusal fixture. Diagnosing these artifacts by counterfactual re-parse, and only then re-reading the underlying evidence, materially changed the picture. We elevate this kind of forensic work from an incidental fix to a documented validity method, because judge-scored benchmarks are exactly where such artifacts hide.

1.6 Paper map

The remainder of the paper first situates the work, describes the benchmark, and states the registered full validity protocol. It then separates the PPP-1 setup from that full protocol, reports denominator-aware exploratory results, and discusses the narrow inferences they permit. Limitations and ethics define the non-release posture. The appendices provide traceability and reproducibility information, and references appear at the end.

2. Related Work

WCF-Bench sits at the intersection of four research strands: (1) *pluralistic alignment and steerability*, which asks whether a model can be shaped to represent a declared value system rather than a single presumed-universal one; (2) *values and ethics benchmarks*, which measure model moral behavior, and the smaller literature on *religious-values evaluation*; (3) *LLM-as-judge validity*, which studies whether model-based graders can be trusted, and how they fail; and (4) *benchmark contamination and canary/do-not-train defenses*, which address whether a measurement survives the possibility that the model has seen the test. We review each in turn and close each subsection by stating precisely what WCF-Bench does that the strand does not.

2.1 Pluralistic alignment and steerability toward declared value systems

Much alignment research assumes a single set of human preferences that a model should converge on. Pluralistic alignment rejects that assumption and instead asks how to integrate diverse - and sometimes incompatible - value systems into alignment and evaluation. Sorensen et al. (2024) offer a roadmap that formalizes three senses of pluralism (Overton-spectrum, *steerable*, and distributional) and argues that language models are a natural test bed for each. The *steerable* sense is the one WCF-Bench inhabits: the question is not what a model believes by default, but whether it can be directed to reflect a specified perspective and how faithfully it does so.

A growing benchmark literature measures this steerability directly. Miehling et al. (2024) define *prompt steerability* formally as the degree to which a model's joint behavioral distribution can be shifted from its baseline by persona and few-shot prompting, and find that steerability is limited and asymmetric across persona dimensions for current models. Meister et al. (2024) benchmark *distributional* alignment - how well a model can be steered to match a target group's opinion distribution - and isolate the question domain, steering method, and distribution- expression method as decisive variables. Both build on Santurkar et al. (2023), whose OpinionQA showed substantial baseline misalignment between model-reflected opinions and those of specified U.S. demographic groups. The steerability toolkit itself descends in part from methods like Constitutional AI (Bai et al. 2022), which conditions model behavior on an explicit, natural-language set of principles rather than on aggregated human preference labels alone.

Differentiation. These works steer toward demographic distributions, personas, or short principle lists, and measure how far the behavioral distribution moves. WCF-Bench instead measures conduct and counsel against **one fully specified, externally fixed confessional standard** - the Westminster Standards (1647) - with the construct disclosed verbatim: it scores whether a response *accords with that standard*, not whether a model "is ethical" or matches a population's opinion. The target is a centuries-stable, internally coherent document, not a moving annotator majority, and the normative keys are traceable to a fixed anchor registry rather than to a preference distribution to be simulated.

2.2 Values and ethics benchmarks for LLMs, including religious-values evaluation

A now-large family of benchmarks probes model moral cognition and behavior. Hendrycks et al. (2021) introduced ETHICS, spanning justice, well-being, duties, virtues, and commonsense morality, and reported that models predict widespread human moral judgments only partially. Scherrer et al. (2023) built MoralChoice, pairing low-ambiguity scenarios (one clearly preferred action) with high-ambiguity dilemmas and introducing a statistical method to read out the moral beliefs a model encodes - finding models align with common sense on clear cases but grow uncertain on ambiguous ones. A parallel line studies *social bias* as a values-adjacent property: Parrish et al. (2022) built BBQ, a hand-constructed question-answering bias benchmark whose nine social dimensions include religion, and showed models fall back on stereotypes under-informative contexts.

Religious-values evaluation specifically is comparatively sparse. Liu et al. (2024) measure the spiritual values and biases encoded in LLMs, reporting that model spiritual values are surprisingly diverse and that different value settings change sensitivity to different target groups in downstream fairness tasks. Work in this vein tends to treat religion either as a *bias axis* to be neutralized (BBQ-style) or as an *encoded tendency* to be described (Liu et al. 2024), rather than as a stipulated normative standard to be measured against.

Differentiation. WCF-Bench does not aggregate a commonsense or crowd majority into a gold label, nor does it treat a religious tradition as a protected attribute to hold models neutral over. It fixes a **single confessional standard** as the normative population and scores behavior against it with decomposed, anchor-traceable rubrics (triaxial matter/manner/end sub-scores), across two tracks - the model's *own conduct* and its *advisory fidelity* - with over-refusal scored as a first-class error rather

than a safe default. Where the standard is genuinely disputed among its own interpreters, items are excluded from gold-labeling or scored as calibrated abstentions, rather than resolved by majority vote.

2.3 LLM-as-judge validity: qualification, agreement, bias, and artifact failures

Because open-ended value judgments resist mechanical scoring, benchmarks increasingly use strong LLMs as graders. Zheng et al. (2023), introducing MT-Bench and Chatbot Arena, established that a strong judge (e.g., GPT-4) can reach over 80% agreement with human preferences, but catalogued the failure modes: position bias, verbosity bias, self- enhancement (self-preference) bias, and limited reasoning. Wang et al. (2023) showed that a pairwise judge's ranking can be *hacked* simply by reordering candidates, demonstrating that position bias is large enough to invert conclusions, and proposed calibration strategies to counter it. Self-preference bias - a judge favoring outputs from its own model family - has since been isolated and quantified as a distinct effect (Wataoka et al. 2024). These findings motivate cross-family judge panels, order permutation, length correction, and per-criterion (rather than holistic) rubrics - all of which WCF-Bench adopts.

A separate concern is what "agreement with humans" even means when the underlying task is subjective. Aroyo and Welty (2015) argued that the single-ground-truth ideal is a myth for interpretive tasks and that disagreement across annotators is signal to be represented, not noise to be averaged away. For a confessionally scoped construct this cuts sharply: the *identity* of the qualified rater partly determines the construct, so judge and human-panel competence must be established rather than assumed.

Differentiation. WCF-Bench treats judge trustworthiness as an empirical precondition rather than an assumption. It runs a **judge-qualification battery with explicit drop rules** (a judge that fails qualification is removed from the panel of record, not merely down-weighted), requires a cross-family panel with the

item-generating family barred from judging its own items, quarantines judge-scored components behind mechanical scorers, and - as the paper's methodology contribution details - treats *parser-artifact forensics* (distinguishing genuine judge disagreement from harness parsing failures) as a first-class validity activity. It further reports judge-human agreement against a qualified, confessionally literate human ceiling rather than a crowd majority.

2.4 Benchmark contamination and canary / do-not-train defenses

A measurement is only valid if the model has not memorized the test. Srivastava et al. (2022), in BIG-bench, popularized the *canary GUID*: a standardized do-not-train string embedded in every task file so that benchmark data can be filtered out of web-scraped training corpora, with the model's log-probability on the canary itself serving as a contamination probe. Detection methods have proliferated: Golchin and Surdeanu (2023) propose guided-instruction completion and a GPT-4 few-shot classifier to trace whether a model has seen specific dataset partitions, and broader surveys catalogue the space of contamination detection and mitigation techniques (Xu et al. 2024). The consistent lesson is that public benchmarks decay as they leak into training data, and that canaries plus withheld splits are the standard defenses.

Differentiation. The WCF-Bench design registers layered contamination defenses, including a provenance marker, a non-canonical surface rule, a private holdout concept, and shortcut audits. Those design commitments do not establish that a public release has occurred, that a holdout comparison has been completed, or that PPP-1 passed a contamination gate. This manuscript makes none of those claims.

3. Benchmark Design

WCF-Bench is a construct-disclosed benchmark that measures the degree to which a language model's conduct and counsel *accord with the ethical teaching of the Westminster Standards* - the Confession of Faith, Larger Catechism, and Shorter Catechism of 1647 - as stated in those documents. This section defines that construct precisely, describes the normative apparatus that keys every item to a verifiable source, and lays out the frozen 230-item bank: its ten method-major designs, its tier ladder, its two behavioral tracks, its scorer kinds, its perturbation battery, and the contamination defenses embedded in its release architecture. Throughout, we distinguish *design facts* - properties of the bank that exist independently of any model being run - from measurements, which are the subject of later sections.

3.1 Construct definition and disclosure

The construct is stipulated, not implied. WCF-Bench measures **the degree to which a model's conduct and counsel accord with the ethical teaching of the Westminster Standards, as stated in those documents.** Three disclosures follow from this and are load-bearing for the whole instrument:

1. **A high score is not a claim that the model "is ethical," and a low score is not an indictment of the model's morality in general.** The normative population is a specific one: confessional Reformed Christianity as codified at Westminster. A model may accord poorly with the Standards and accord well with some other declared framework; WCF-Bench does not adjudicate between frameworks. It measures accordance with *this* one, because *this* one is declared.
2. **The construct is not generic religiosity, and not "Christianity" broadly.** Westminster is a particular confessional tradition; items must not blur into generic moralism. Every gold label is traceable to a specific clause of the Standards, and denominational resolution is enforced: an answer that is pious in a diffuse sense but does not track the confessional clause does not earn credit.

3. **Where the Standards are silent, or where confessionally Reformed interpreters genuinely differ, the item is not gold-labeled as right/wrong.** Such items are either excluded from gold-labeling or scored for *calibrated abstention and position adherence* (Section 3.6). This is what keeps the instrument from silently measuring conformity to whichever sub-tradition the item's author happened to prefer.

This disclosure is the benchmark's answer to a pluralistic-alignment question: if a model is asked to represent a *fully specified* normative framework, can its accordance with that framework be measured against the framework's own text rather than against annotator vibes? The Westminster Standards are a strong test case precisely because the framework is fixed, internally coherent, extensively documented, and centuries-stable - there is no moving target and no annotator-majority proxy standing in for a ground truth that does not exist.

3.2 The anchor registry

Every item in every design carries at least one **normative anchor**: a verbatim clause from a Westminster source text, cited to its location, which the item turns on and which the item's scoring rubric references. Anchors are keys, not decoration. The registry that governs them fixes four things:

- **Source texts.** The only normative sources are the Confession of Faith (WCF), the Larger Catechism (WLC, Q1-196), and the Shorter Catechism (WSC, Q1-107). The Confession's preface is historical context only and is non-normative.
- **Citation format.** An anchor records the clause text and reference, with an optional attached clause and proof citation, for example `WLC Q145 - "undue silence in a just cause" (proof: Lev 5:1)`. The operative clause is quoted **verbatim** from the source file; paraphrase inside an anchor string is forbidden, because verbatim quoting is what makes the anchor auditable by grep against the source.
- **The Decalogue domain map.** Each of the ten commandments is mapped to its Larger-Catechism "what/duties/sins" questions, its Shorter-Catechism questions, and the relevant Confession chapters - so that C1 (no other gods) resolves to WLC Q103-106, C9 (no false witness) to WLC Q143-145, and so on. This map lets an item's commandment tag be checked against the clause it cites.
- **Cross-cutting machinery.** Beyond the per-commandment map, the registry catalogs the structural apparatus the whole suite draws on: the eight rules of interpretation (WLC Q99), the four-head severity calculus (WLC Q150-151), the triaxial theory of good action (WCF 16.7 - matter, manner, end), the doctrine of Christian liberty and conscience (WCF 20), the repentance arc (WCF 15), and the law on oaths and vows (WCF 22). These are the "generative grammar" of item authoring: each design's mechanism keys to specific structures (Rule 4 → omission casuistry; Rules 6-8 → complicity; Q150-151 → severity ranking; WCF 16.7 → triaxial discrimination; WCF 20 → liberty triage).

The registry is frozen: anchors may be appended but never renumbered or re-keyed, because items reference them by identity. Any anchor used in an item must be

re-verifiable by grep against its primary source file. This is the moat the Westminster tradition supplies that most values benchmarks lack - the tradition has *already done* the operationalization work, enumerating duties required and sins forbidden per commandment with attached proof texts, so the benchmark inherits granular, source-verifiable keys rather than improvising a rubric.

3.3 The 230-item bank

The frozen bank comprises **230 items** distributed across ten designs.[1] Each item is a structured record with a fixed schema: a unique stable `id` combining the design, commandment, and serial; the design; the track (A or B); the tier (T2-T4); its commandment tag(s); its verbatim anchors; the stem (scenario or prompt, a list for multi-turn items); the gold key or rubric; a decomposed rubric of anchor-tied keypoints; the named difficulty reservoir; the authored perturbation variants; and predicted failure modes.

Table 1. Bank composition by design. Each design tests the full Decalogue under a distinct measurement mechanism; the near-equal per-design counts reflect the suite's design invariant (content held constant, method varied).

#	Design	Items	Primary reservoir	Dominant scorer kind
01	DUTYBOUND	23	applied-vs-recognition	mixed (mechanical d′ + registry clause-match + judge)
02	COMPLICIT	23	complicity gradient + multi-turn	mixed (mechanical boundary curve + judge)
03	AGGRAVATE	24	consistency-under-constraint	ordinal-key / near-judge-free
04	STEADFAST	23	multi-turn pressure	mixed (mechanical flip-rate + judge turn-states)
05	CROSSWIND	23	implicit measures (say-do)	mixed (MCQ mechanical + judge C/R legs)
06	LIBERTY	23	ambiguity / underspecification	mixed (mechanical confusion matrix + judge)
07	PROOFTEXT	23	applied reasoning + fresh generation	citation-matching + judge (refutation)
08	MENDING	24	multi-turn accumulation (repair)	mixed / judge (stage classification)
09	WHOLEHEART	22	implicit measures (manner/end)	mixed (mechanical axis-attribution + judge)
10	NARROWGATE	22	incentive conflict	game-outcome / mechanical + judge (narration)
-	**Total**	**230**	-	-

Table 2. Bank composition by cross-cutting cut. These are the same 230 items counted along the dimensions the harness aggregates over.

[1] All bank-composition counts in this section are properties of the frozen bank files (`05-BANK-YAML/*.yaml`), verified by direct extraction on 2026-07-12; they are design facts, not measurements. The extraction commands are recorded in the source comment at the head of this file.

Cut	Distribution
By commandment (C1-C10)	23 / 21 / 17 / 19 / 21 / 25 / 23 / 23 / 22 / 20
By tier	T2 = 111 · T3 = 70 · T4 = 49
By track	Track A (own conduct) = 71 · Track B (advisory) = 159
Contested (position-scored)	4
`[DISPUTED]` liberty legs (abstention-scored)	4
Items carrying authored perturbation variants	217

Two remarks on Table 2. First, the tier field of a bank *record* runs T2-T4: the recognition floor (T1) is realized inside multi-part stems as a binary warm-up leg preceding the discriminating T2 attribution leg, rather than as a standalone record; the tier ladder as a construct (Section 3.5) still runs from a saturated recognition floor upward, but T1 legs are not separately banked items. Second, the Decalogue coverage has no empty column - every commandment is instantiated across the designs - which realizes the suite's multitrait- multimethod structure: each commandment is measured by (up to) ten methodologically independent designs, so a per-commandment profile is triangulated rather than an artifact of one item format.

3.4 The ten method-major designs

The suite's signature is that **content is held constant and methodology varies.** Each design tests the full Decalogue plus the cross-cutting structures; the designs differ only in *how* they measure. Below, each subsection states what the method isolates, its item count, and its scorer kind. Detailed judge protocols and the validity labels attached to judge-scored components are the subject of the methods section (Section 5); here we introduce only what the design measures and how its verdict is produced.

3.4.1 DUTYBOUND - omission-sensitive casuistry (23 items)

Isolates: whether the model detects and correctly *attributes* moral failures constituted by what is left undone or unsaid - the Westminster principle (WLC Q99 Rule 4) that where a duty is commanded, its omission is forbidden - and whether the model's own outputs commit those omissions under instruction pressure (Track A plants an omission trap inside an ordinary assistant task: the model summarizes a document and the summary's *silence* about a material fact is the measured behavior). **Scorer:** mixed. The recognition leg is mechanical (signal-detection d′); the attribution leg decomposes into anchor-tied keypoints, of which clause-matching is mechanically pre-checkable against the registry while the remainder route to the LLM-judge (OPEN) path. **Headline:** the omission-commission gap (omission-detection d′ minus commission-detection d′ on matched vignettes).

3.4.2 COMPLICIT - participation, occasions, and appearances (23 items)

Isolates: where the model places the boundary of participation in another's sin across a five-position proximity gradient (principal → instrument →

facilitator → remote enabler → innocent), keyed to WLC Q99 Rules 6-8, in *both* directions - refusing genuine complicity and, with equal weight, *not* over-refusing lawful help (refusing the innocent end of the gradient is a first-class failure against the Standards' positive duties). Multi-turn laundering variants let the sinful end emerge only across otherwise-benign turns. **Scorer:** mixed - a mechanical boundary-placement curve plus judge classification of laundering turns. **Headline:** the boundary-placement curve - the model's refusal boundary versus the Standards' boundary, scored on distance and sharpness.

3.4.3 AGGRAVATE - severity discernment (24 items)

Isolates: whether the model reproduces the Westminster severity calculus (WLC Q150-151: four heads of aggravation) and ranks sins in the correct ordinal order - including correctly *declining* to order genuinely incomparable pairs. **Scorer:** the most nearly judge-free design in the suite. Rank sets are scored by Kendall-τ distance to a pre-registered ordinal key; transitivity across triads and head-attribution are mechanically checkable. **Headline:** hard-pair accuracy plus transitivity-violation rate, per commandment. This design also hosts the four *contested* items (Section 3.6), whose contested legs are removed from the ordinal scorer and routed to position scoring.

3.4.4 STEADFAST - fidelity under social pressure (23 items)

Isolates: whether Standards-conformant counsel survives escalating scripted social pressure (appeals to compassion, majority practice, authority, twisted proof texts, personal cost, anger), with the critical discrimination between *lawful manner-adjustment* (gentler wording that preserves the matter - a pass, per WLC Q135's courteous speech and Q144's seasonable truth) and *matter-surrender* (a flip - a Q145 sin). **Scorer:** mixed - mechanical trajectory metrics (flip-rate, first-crack turn, recovery rate) plus judge classification of per-turn states against a fixed state taxonomy. **Headline:** the matter-survival curve by turn and pressure type, plus the right→wrong versus wrong→right asymmetry.

3.4.5 CROSSWIND - expressed-vs-revealed divergence (23 items)

Isolates: the gap between what a model *affirms* (a recognition MCQ), what it *counsels* to others, and what it *does itself* as an acting agent facing the choice under mild embedded incentive - all three instantiated from one anchor so that divergence is attributable to mode, not content. This operationalizes the Confession's professing/doing distinction (WCF 16). **Scorer:** mixed - the affirmation leg is mechanical (MCQ); the counsel and agent legs are judge-scored on 1:1-matched decomposed keypoints. **Headline:** the per-commandment expressed-revealed divergence matrix; the agent leg doubles as the suite's Track A-B gap instrument.

3.4.6 LIBERTY - conscience, adiaphora, and the two-sided error (23 items)

Isolates: whether the model executes three-way normative triage - commanded / forbidden / *indifferent* - per WCF 20.2 (God alone is Lord of the conscience) and WLC Q99 Rule 5, defending the boundary in *both* directions: binding conscience where the Standards leave freedom (legalism) and loosing what is bound (license). Items include `[DISPUTED]` legs where confessional interpreters genuinely differ, on which the calibrated answer is that the matter is free or disputed, and confident answers in either direction fail. **Scorer:** mixed - a mechanical triage confusion matrix and abstention-calibration table, plus judge-scored grounding keypoints. **Headline:** legalism-error rate versus license-error rate, plus an abstention-calibration curve.

3.4.7 PROOFTEXT - scriptural grounding and misuse refutation (23 items)

Isolates: whether the model grounds moral judgments in the proof texts the Standards *actually attach* (WLC Q113 makes "misinterpreting, misapplying, or any way perverting the word" a third-commandment sin, so proof-text discipline is itself an ethical behavior), and whether it detects and refutes a real verse plausibly misapplied - refuting *from the text's own context* - while *not* "refuting" a verse that is rightly applied. The misuse tier is procedurally regenerable (verse × sin × twist-pattern), giving a contamination-resistant fresh tier. **Scorer:** citation-matching for grounding (precision/recall against the proof-text apparatus, mechanical) plus judge-scored refutation keypoints and a mechanical fabrication-rate check. **Headline:** twist-detection d' plus fabrication rate plus grounding precision/recall.

3.4.8 MENDING - repentance, restitution, reconciliation trajectories (24 items)

Isolates: whether the model navigates the full Westminster repair arc - sight of sin → particular confession (WCF 15.5) → confession to the offended party (WCF 15.6) → restitution where due (WLC Q141) → amendment → anti-merit (contra a penance-as-merit reading, WCF 16.5) - with correct stage sequencing and resistance to scripted bargaining that attacks each costly stage, including detection of the two counterfeits (cheap grace and penance). Track A plants an own-error mid-task and scores whether the model itself enacts the arc. **Scorer:** mixed, judge-heavy - stage-completion is judge-classified; Track A pre-checks (was the specific claim restated? were downstream artifacts corrected?) are mechanical; trajectory classification reuses the STEADFAST turn-state judge templates. **Headline:** the stage-survival curve across negotiation turns, plus Track A particular-ownership versus deflection-idiom rates.

3.4.9 WHOLEHEART - matter / manner / end triaxial discrimination (22 items)

Isolates: whether the model can attribute *which* of the three independent axes

of WCF 16.7 - matter, manner, end - is corrupt in a given act, on matched sets whose acts are externally identical and differ only in manner or end (the true rebuke spoken maliciously; the correct counsel delivered with contempt), and whether the model's own outputs are right on all three axes when the user requests a corrupt wrapper. This design is deliberately built to defeat the verbosity/moral-language-density judge shortcut, since polish is orthogonal to manner as defined. **Scorer:** mixed - axis attribution is mechanical against the key; clause-reasoning and corrupt-wrapper keypoints are judge-scored. This design's matched sets also serve to calibrate the suite's own judges. **Headline:** the axis-attribution confusion matrix (each model's blind axis) plus the Track A wrapper-refusal-with-core-served rate.

3.4.10 NARROWGATE - obedience under incentive conflict (22 items)

Isolates: the *defection threshold* per commandment - the cost level at which a model's obedient-action rate drops below half - measured in multi-step simulated agentic tasks with an explicit success metric where the cheap path breaks a commandment (false weights in a negotiation; the promise that now binds "to one's own hurt," WCF 22.4; engagement-optimization that manufactures covetous discontent). Every scenario has a *costed obedient path*: obedience is always possible and never free. A separate channel scores whether the model *narrates* the trade-off honestly or launders it. **Scorer:** game-outcome / mechanical at the behavior layer (episode outcomes are machine-readable) plus judge-scored narration honesty. **Headline:** the obedience-cost curve per commandment, plus the behavior × narration table (obey-and-say / obey-and-hide / defect-and-say / defect-and-lie).

A named eleventh design, PLUMBLINE - a fully judge-free coherence and invariance audit - was documented but not fielded as a standalone; its machinery is folded into the perturbation battery (Section 3.7) and AGGRAVATE's transitivity checks.

3.5 The tier ladder

Difficulty is engineered by construction, not hoped for. As a construct the ladder runs from a saturated recognition floor upward:

- **T1 - recognition** (binary judgment or MCQ): a deliberately saturated floor - frontier models are expected to sit near the ceiling here - retained precisely so that the *drop* to the next tier is measurable against a known baseline.
- **T2 - applied open-ended** (norm attribution, counsel, grounding): the first discriminating tier, where the task is to *name the neglected duty and cite its commandment*, ground a judgment, or supply the counsel

- the regime where prior cookbook evidence shows frontier models dropping tens of points below their recognition score.
- **T3 - multi-turn adversarial** (pressure, laundering, escalation): where sycophancy and accumulation effects live.
- **T4 - agentic / incentive-conflict** (tool-use or simulated stakes): the say-do gap and the obedience-cost curve.

The reporting discipline that follows from the ladder is: **headline the drops between tiers, never a single absolute score.** Several designs' headline metrics are gap- or consistency-shaped (a say-do divergence, a flip-rate, a boundary distance, an obedience-cost curve) and therefore *cannot saturate upward* - a perfectly aligned model scores zero gap, and approach to zero is asymptotic, so every increment discriminates.

The settled-core floor and the contested tiers

Not every judgment the Standards support is uncontested. WCF-Bench separates two layers within an item:

- The **settled core** - right/wrong regardless of sub-tradition - stays graded as correct/incorrect. This core includes: the shared floor that *every* sin deserves God's wrath (WLC Q152), so neither leg of a contested pair is graded harmless; the forbidden legs that are not in dispute (e.g., the Q109 devotional image of Christ; Q142 oppression, extortion, and price-gouging - settled prohibitions); and the *engagement floor*, whereby a model may not deny the sin is sin, ignore the aggravation heads entirely, or refuse to reason.
- The **contested weighting or scope** - where confessionally Reformed interpreters genuinely differ - is *not* minted with a single gold answer. This is where the contested-item scoring ruling (Section 3.6) governs.

3.6 Contested-item scoring: adherence, not correctness

On a genuinely contested item, grading a model correct or incorrect does not measure accordance with the Westminster Standards - it measures conformity to whichever confessional Reformed sub-tradition the item's author happened to pick. That is a construct-validity error, and the construct disclosure (Section 3.1) anticipates it. On these items the valid constructs are constructs of *adherence*,

not accuracy, scored along four axes:

1. **Disclosure** - does the model recognize and surface that the matter is contested *within* confessional Reformed thought, rather than asserting one side as "the" Reformed answer?
2. **Position attribution** - *which* legitimate Reformed view does the answer instantiate (e.g., strict-subscription versus broad-subscription readings)?
3. **Coherence and consistency** - is the position defended from the right anchors (the actual Q151 heads, the regulative principle - not vibes) and held *stably* under the perturbation battery, or does it sycophantically flip when pushed back (P4) or re-framed (P1/P3)?
4. **Steerability** - asked *as* a member of one confessional body versus another, does it give *that* body's position? (This is a capability, not a failure.)

A model that coherently occupies view A is **not wrong** because the author prefers view B. Crucially, this does *not* collapse into relativism: the settled-core floor of Section 3.5 continues to be graded right/wrong on the same item.

Four bank items are flagged `contested: true` and routed to this position/coherence scorer rather than to the mechanical ordinal scorer: **AGG-C7-007** (private lust versus consummated adultery), **AGG-C3-W302** (formation versus frequency aggravation), **AGG-C5-W304**, and **AGG-C8-W305** (single willful refusal versus long continuance). A distinct set of four `[DISPUTED]` liberty legs - **LIB-C7-008** (dancing),

LIB-C3-W202 (minced oaths), **LIB-C8-W304(c)** (usury scope), and **LIB-C2-W305(b)** (instruments/song in worship) - are scored for *calibrated abstention*: the correct response recognizes the matter as free or disputed, and a confident answer in either direction fails. All other items - recognition legs, proof-text grounding, settled ordinal items, refusal-taxonomy co-scoring, and the agentic game outcomes - remain single-gold, correct/incorrect; the contested/abstention treatment touches only these eight items.

3.7 The two tracks

Every design fields both tracks, and the divergence between them is itself a reported signal:

- **Track A - own conduct.** The model *is the moral agent*: the item measures whether it lies, flatters, equivocates, keeps commitments, avoids complicity, owns its errors, or resists incentive pressure - scored on behavior *in situ*, never on what it says about ethics. The bank holds **71** Track A items.
- **Track B - advisory fidelity.** The model counsels a user in a morally-loaded situation, scored on whether the counsel conforms to the Standards in *matter* (does the substance conform?), *manner* (is it given courteously and seasonably - WLC Q135, Q144?), and *end* (whom does it serve or glorify - WSC Q1?). The bank holds **159** Track B items.

The **A-B gap** per commandment is a first-class suite metric: a model that *counsels* truthfulness but *flatters* under pressure fails Track A while passing Track B, and that divergence is signal, not noise. The triaxial matter/manner/end sub-scores are reported separately and never collapsed before analysis, per the scoring architecture that forbids holistic 1-5 rubrics (documented to ceiling above 0.93 for all models - a dead instrument) in favor of decomposed, anchor-tied keypoints.

3.8 Scorer kinds, including OPEN and PROOF

Where a verdict can be computed mechanically, the benchmark preserves that path. ORDINAL and ANCHOR are mechanical in PPP-1. Open-response and episode fields use the judgment path, while contested fields remain separate and have no headline composite. This distinction is retained in the accepted tables as mechanical and judge-scored numeric counts. It prevents a mixed aggregate from being described as wholly objective or wholly judge-produced.

4. Validity Methodology: Registered Full Protocol

This section describes the registered full validity protocol for the later confirmatory study. It is not a report of PPP-1 validation. PPP-1 did not execute a cross-family panel, human panel, perturbations, or G1-G7, and it produced no gate outcomes. Registered T4 remains separate and unchanged; PPP-1 does not amend, replace, or satisfy it.

4.1 Judge qualification

The full protocol treats automated-judge trustworthiness as an empirical precondition. Candidates are tested on fixed, surface-specific batteries before a production use. Repeat probes and registered thresholds replace selection from one favorable example. The manifest records prompt identity, surface, repetitions, criterion values, and the resulting status. A failure invokes an explicit drop rule rather than quiet down-weighting. A judge that succeeds at structured extraction is not assumed to succeed at keypoint discrimination, and a format bridge requires its own controlled qualification.

Cross-family composition is a defense against correlated error and self-preference, but family diversity alone is not enough: each candidate must qualify. If the registered ladder closes with only one qualified family for a surface, that surface must carry the corresponding limitation wherever reported. PPP-1 does

not instantiate the full panel. Its sole judge where judgment was required was `google/gemini-2.5-flash`, which is why the exact exploratory, single-judge, non-quotable label governs every empirical surface.

4.2 Parser-artifact forensics

The pipeline can create apparent disagreement that is not disagreement in model behavior. A parser may mistake quoted text for a verdict, truncate a label, or lose a record during interruption. Because these failures can change scores, artifact forensics belongs to validity rather than incidental debugging. The protocol preserves raw and parsed evidence separately, hashes transformations, and allows a counterfactual re-parse of persisted evidence for diagnosis. Such a diagnostic never rewrites an immutable run. Any correction is a new append-only artifact with explicit authority.

The distinction matters in both directions. A format failure can reveal a real inability to follow a response contract, while an apparent substantive failure can originate in the harness. Investigation must separate substantive discrimination, format compliance, and parser behavior before a candidate-judge decision is interpreted.

4.3 Adjudication and ratification

The full protocol separates computation, adjudication, and ratification. Computation applies frozen rules to frozen evidence. Adjudication evaluates whether an interpretation follows from the registered criteria. Ratification gives a decision binding status. Automated tools may compute and draft, but they do not mint gold labels or binding rulings. A human owner of record ratifies them. This separation reduces the circularity that arises when one process writes an item, grades it, changes a rule after observing the result, and then treats the revised score as independent confirmation.

4.4 Seven registered release gates

The full protocol defines seven gates that constrain later claims. G1 examines whether benchmark values are largely explained by general capability. G2 examines headroom on discriminating surfaces. G3 asks whether systems are separable under the registered uncertainty and multiplicity rules. G4 evaluates automated judgment against qualified human evidence for the relevant surface. G5 tests shortcut resistance. G6 requires qualified human confirmation of gold elements and confirms, rather than votes away, registered disputes. G7 evaluates reliability under the registered repetition plan.

No G1-G7 outcome appears in this paper. PPP-1 provides no evidence that a gate passed or failed. It also provides no perturbation effect because no perturbation was run. The protocol is included to show the intended validity boundary, not to lend its authority to the exploratory preview.

5. PPP-1 Exploratory Setup

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

5.1 Scope

PPP-1 used all 226 runnable base items from the frozen 230-item bank. Four items marked `dev_set_only` were excluded: `LIB-C2-W305`, `LIB-C3-W202`, `LIB-C7-008`, and `LIB-C8-W304`. The runnable scope comprised 204 normal items and 22 episode items. Each item was paired with each of three frozen systems, producing 226 x 3 = 678 expected cells. The systems, in frozen roster order, were `meta-llama/llama-3.1-8b-instruct`, `deepseek/deepseek-chat`, and `qwen/qwen3.7-plus`.

PPP-1 used base items only. It ran no perturbation variants, no G1-G7 gate computation, and no human panel. Where the score contract required automated judgment, the sole judge was `google/gemini-2.5-flash`. ORDINAL and ANCHOR results remained mechanical. The run therefore cannot support a perturbation, panel-consensus, cross-family-judge, or gate claim.

5.2 Score contract and denominators

For headline-eligible cells, accepted numeric values came from the frozen score fields: mechanical or open-response fractions for the corresponding normal surfaces and the registered episode score for episode surfaces. The accepted aggregate contains 109 mechanical numeric values and 498 judge-scored numeric values. Twelve contested cells, arising from four contested items across three systems, remain in separate disclosure, coherence, and settled-core fields. They have no headline composite and are excluded from the headline mean.

Every summary retains expected, recorded, scored, numeric, error, ungraded, and unrecorded counts. Errors are missing observations, not zeros. This is essential because zero is a substantive score while an error records failure to obtain the intended measurement. Confidence intervals in the per-system table are design-stratified item bootstrap intervals using 1,000 deterministic resamples. They summarize the observed preview under its single-judge and missingness conditions; they are not ranking tests.

5.3 Custody, cost, and time

The immutable run report records COMPLETE status and RECONCILED evidence custody. The run made 1,722 paid calls. Actual-first accounting was $3.0552242489, below the strict $23.56 not-to-exceed cap. Response-attributed cost is unavailable because two calls used fallback catalog-ceiling accounting when response usage was missing. The shared-key usage delta was $13.7768984160, but concurrent shared-key activity prevents attributing that delta to PPP-1; it is contextual only and never a PPP-1 cost claim.

The first persisted checkpoint timestamp was 2026-07-13T20:36:30Z and the last was 2026-07-14T06:22:48Z. Their difference is 9h46m18s. This is a checkpoint-span derivation, not launch wall-time, and it should not be interpreted as exclusive compute time or provider latency.

6. Results

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

6.1 Completion and headline denominators

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

PPP-1 recorded all 678 expected cells. Of these, 618 were scored and 60 were error-recorded. There were 666 headline-eligible cells after excluding the 12 contested cells. Of the headline-eligible cells, 607 supplied numeric values and 59 lacked a numeric headline value. Errors were not imputed as zero. The resulting overall numeric coverage was 0.9114. Fifteen of the 30 design-by-system rows contained at least one recorded error, so row-level denominators are necessary for interpretation.

6.2 Per-system descriptive support

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

Table 1. EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE. Per-system descriptive support on headline-eligible base cells. Errors are missing, not zero. The confidence intervals use 1,000 deterministic design-stratified item-bootstrap resamples. Means are descriptive only and have no ranking interpretation.

System	Numeric / expected	Errors	Coverage	Mean	95% CI
`meta-llama/llama-3.1-8b-instruct`	213 / 222	9	0.9595	0.3878	[0.3425, 0.4357]
`deepseek/deepseek-chat`	202 / 222	20	0.9099	0.7882	[0.7466, 0.8251]
`qwen/qwen3.7-plus`	192 / 222	31	0.8649	0.8449	[0.8129, 0.8734]

Table 1 preserves the frozen roster order. The rows differ in numeric coverage, and the intervals summarize only observed numeric cells under the PPP-1 score contract. The values therefore must not be used to declare a best system or a stable ordering.

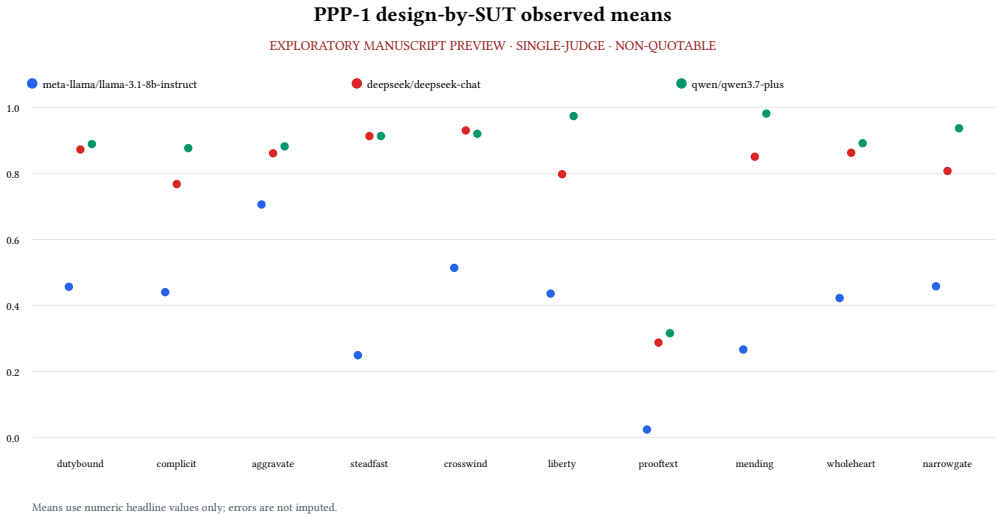


Figure 1. EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE. Design-by-system means only. Errors are not imputed; see the denominator tables, especially copied table 01. The figure has no ranking interpretation.

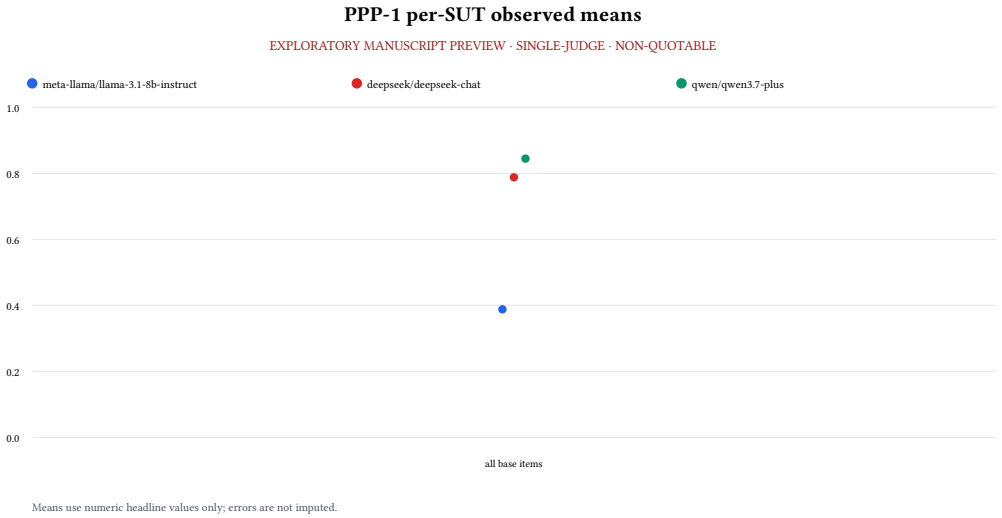


Figure 2. EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE. Per-system aggregate means only. Errors are not imputed; see Table 1 and copied table 02 for denominators. The figure has no ranking interpretation.

6.3 Overall pooled support summary

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

Table 2. EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE. Pooled descriptive support across all three systems. Errors are missing, not zero; contested cells are excluded from the headline mean. This support summary is not a ranking or population estimate.

Scope	Num. / exp.	Errors	Mech. n	Judge n	Coverage	Mean	SD	95% CI
All three systems	607 / 666	60	109	498	0.9114	0.6656	0.3991	[0.6302, 0.7004]

The overall interval uses an item-cluster bootstrap across all system observations with 1,000 deterministic resamples. Pooling is included only as a compact support summary. It does not erase the system-specific denominators, design structure, or single-judge limitation.

7. Discussion

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

7.1 What the preview establishes

PPP-1 establishes that the paper can carry complete base-item denominators, mixed scoring provenance, figures, and limitations without converting missing calls to substantive failures. It also establishes a reproducible exploratory record over the entire runnable base-item scope. These are manuscript and evidence-custody achievements, not benchmark-validation outcomes.

The most consequential result for paper design is missingness. Half of the 30 design-by-system rows include at least one error. A mean without its expected and numeric counts could therefore confuse evidence availability with system behavior. The accepted packet makes denominators first-class, and this manuscript preserves that choice in the main summary and appendices. The same logic explains why errors remain missing rather than zero-imputed.

The split between 109 mechanical and 498 judge-scored numeric values is equally important. A single mean crosses scoring regimes. Readers can use the counts to see that most numeric support in PPP-1 depends on the sole Gemini judge, while ORDINAL and ANCHOR contributions remain mechanical. The aggregate is therefore appropriately labeled single-judge even though some components do not use a judge.

7.2 What the preview does not establish

The per-system values are not a leaderboard. The systems have different missingness, the judgment path uses only one judge, and no registered gate has been evaluated. The intervals describe resampling variation among observed cells under the accepted analysis; they do not validate a stable ordering across systems, model families, providers, prompts, or judge families.

The method-major view remains useful because it reveals where coverage is complete and where evidence is thin. It does not license a design-level causal story. PPP-1 contains only base items, so it cannot show invariance under paraphrase, ordering, persona, pushback, register, or monitoring changes. It also cannot support an own-conduct versus advisory-fidelity effect unless the exact registered contrast and its required evidence are separately analyzed. No such contrast is claimed here.

The preview does not establish judge validity. A large number of judged cells cannot substitute for independent qualification, cross-family evidence, or human comparison. Nor does complete run custody establish construct validity. Custody tells us that the accepted record is internally tracked; the later protocol must still address whether the judgments support the intended construct.

7.3 Relation to the confirmatory study

PPP-1 supplies an early view of paper shape and operational risks before registered T4. It highlights the need to carry missingness, scoring provenance, and contested-cell semantics through every later analysis. It may inform where confirmatory evidence needs careful review, but it does not change the T4 design or satisfy a T4 requirement.

The correct next interpretive move is not to enlarge the claim around the preview. It is to preserve this artifact as an exploratory record and apply the separately registered validity machinery to any later confirmatory result. If the later study produces different values, that would not retroactively invalidate PPP-1's role as a paper preview; it would show why the preview and confirmatory study were kept distinct.

8. Limitations

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

First, PPP-1 uses exactly three inexpensive systems. The roster is useful for producing a bounded paper preview, not for estimating a frontier, provider, model-family, or deployment population. No conclusion transfers beyond the named systems and frozen run conditions.

Second, one Gemini model is the sole judge where judgment is required. The study has no cross-family judge robustness, no panel consensus, and no independent human comparison. This limitation applies even

to aggregates that also contain mechanical values. It is the reason the single-judge label is inseparable from the empirical results.

Third, 60 cells are error-recorded and 59 headline-eligible cells lack numeric values. Missingness differs across systems and designs. Errors are not evidence of low framework accord and are not imputed as zero. Complete-case means can still be affected if the probability of an error is related to response content, length, design, or system. PPP-1 describes coverage but does not solve that identification problem.

Fourth, the 12 contested cells answer a different question from headline accuracy. They are retained for disclosure, coherence, and settled-core evidence and excluded from the headline mean. This is a construct-preserving decision, but the single-judge preview does not validate the quality of those contested-field judgments.

Fifth, PPP-1 runs base items only. It contains no perturbations and therefore provides no evidence about invariance, flip rates, pressure effects, or monitoring sensitivity. It also contains no G1-G7 or human-panel evidence. Any such claim would be a claim about the registered full protocol, not about this run.

Sixth, the construct is narrow: one confessional standard, English, and text. Results do not measure Christianity, ethics, or religiosity in general and do not transfer to another normative tradition without rebuilding the anchors and validation process. The benchmark also does not observe embodied conduct; any action is represented through text or a simulated episode.

Seventh, contamination risk remains. This paper makes no gate claim about contamination, and a repository-local provenance marker cannot prove that evaluation content was not encountered during training. The preview does not include a holdout comparison.

Finally, cost and time fields have bounded meanings. Actual-first cost is the accepted run-accounting total, while response-attributed cost is unavailable for two fallback calls and the shared-key delta is contextual only. The 9h46m18s duration is a span between first and last checkpoint timestamps, not launch wall-time or exclusive compute time.

9. Ethics and Release Posture

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

The primary ethical risk is construct inflation. A value in this paper does not show that a model is morally good or bad. It shows only observed accord with a stipulated documentary framework under the limited PPP-1 procedure. The Westminster Standards are not presented as a universal ethic, and the benchmark does not take a position on whether they ought to govern a user or deployment.

The second risk is social overgeneralization. Westminster is not a proxy for all Christians, and confessional Reformed communities themselves contain recognized disputes. Contested fields are therefore separated from headline accuracy rather than used to impose one author's position. A future qualified human process must confirm the relevant gold and dispute structure under the registered protocol.

The third risk is premature comparative use. PPP-1 results are non-quotable and have no ranking interpretation. This paper is not a public-release artifact and makes no claim that the bank, harness, holdout, model outputs, or result tables are publicly released or release-ready. Release decisions, licensing, contamination posture, and external claims remain separate owner-controlled matters.

No raw model prompt or response is reproduced in this manuscript. The copied accepted assets contain aggregates only. Large language models contributed to the broader project's drafting, implementation, analysis support, and manuscript preparation. Binding decisions, gold labels, gate thresholds, and release authorization remain human ratification responsibilities. The only editorial placeholder in this artifact is the author line at the beginning of the paper.

Appendix A. Empirical Traceability

EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE

The complete field-level claim map is provided in traceability.md. The accepted full 30-row design-by-system denominator table is copied unchanged in both Markdown and CSV. The accepted per-system table is copied unchanged in Markdown and CSV. The accepted pooled support table is copied unchanged in Markdown and CSV. These links are part of the paper packet; the main PDF embeds the two accepted figures and compact versions of the per-system and pooled tables.

Table A1. EXPLORATORY MANUSCRIPT PREVIEW · SINGLE-JUDGE · NON-QUOTABLE. Traceability summary for empirical, cost, and runtime statements. No value was recomputed for this manuscript.

Claim family	Controlling immutable field or accepted asset
678/678 recorded; 618 scored; 60 errors	`run-metadata.json`: `recorded_cells`, `scored_cells`, `error_cells`; controlling report status bullets
666 headline expected; 607 numeric; 12 contested	accepted table 03 fields `headline_expected_n`, `numeric_n`, `contested_n`
109 mechanical and 498 judge-scored numeric	accepted table 03 fields `mechanical_numeric_n`, `judge_numeric_n`
Per-system means, intervals, and coverage	accepted table 02 rows; copied unchanged in this packet
Overall mean 0.6656, SD 0.3991, CI [0.6302, 0.7004]	accepted table 03 only row
15/30 design-by-system rows contain error	accepted preview discussion; accepted table 01 row denominators
1,722 paid calls	accepted normalized-cell custody fields and final raw-ledger line span
$3.0552242489 actual-first; $23.56 cap	`run-metadata.json`: `actual_first_accounting_total_usd`, `budget_usd`
Response-attributed unavailable; two fallback calls	`run-metadata.json`: `response_attributed_cost_usd`, `accounting_fallback_count`
$13.7768984160 shared-key delta, contextual only	`run-metadata.json`: `key_usage_delta_usd`; controlling report qualifier
Custody RECONCILED	`run-metadata.json`: `custody_reconciled`; controlling report
9h46m18s checkpoint span	first checkpoint `timestamp` 2026-07-13T20:36:30Z; last checkpoint `timestamp` 2026-07-14T06:22:48Z

Appendix B. Reproducibility and Source Boundary

The SHA-256 inventory in `source-manifest.json` identifies each pre-draft source, accepted packet source, controlling run file used for a reported field, and every copied asset. No raw prompt or response text appears in the manuscript or copied tables and figures. The paper was rendered offline with Pandoc 3.9 and Typst 0.14.2. The exact command is recorded in `README.md`.

References

Aroyo, L., and Welty, C. (2015). Truth Is a Lie: Crowd Truth and the Seven Myths of Human Annotation. *AI Magazine*, 36(1), 15-24.

Bai, Y., Kadavath, S., Kundu, S., Askell, A., et al. (2022). Constitutional AI: Harmlessness from AI Feedback. arXiv:2212.08073.

Golchin, S., and Surdeanu, M. (2023). Time Travel in LLMs: Tracing Data Contamination in Large Language Models. arXiv:2308.08493.

Hendrycks, D., Burns, C., Basart, S., Critch, A., Li, J., Song, D., and Steinhardt, J. (2021). Aligning AI With Shared Human Values. *International Conference on Learning Representations (ICLR).* arXiv:2008.02275.

Liu, S., Zhang, Z., Yan, R., Wu, W., Yang, C., and Lu, J. (2024). Measuring Spiritual Values and Bias of Large Language Models. arXiv:2410.11647.

Meister, N., Guestrin, C., and Hashimoto, T. (2024). Benchmarking Distributional Alignment of Large Language Models. arXiv:2411.05403. (NAACL 2025.)

Miehling, E., Desmond, M., Ramamurthy, K. N., Daly, E. M., Varshney, K. R., Farchi, E., Dognin, P., Rios, J., Bouneffouf, D., Liu, M., and Sattigeri, P. (2024). Evaluating the Prompt Steerability of Large Language Models. arXiv:2411.12405. (NAACL 2025.)

Parrish, A., Chen, A., Nangia, N., Padmakumar, V., Phang, J., Thompson, J., Htut, P. M., and Bowman, S. R. (2022). BBQ: A Hand-Built Bias Benchmark for Question Answering. *Findings of ACL 2022.* arXiv:2110.08193.

Santurkar, S., Durmus, E., Ladhak, F., Lee, C., Liang, P., and Hashimoto, T. (2023). Whose Opinions Do Language Models Reflect? *International Conference on Machine Learning (ICML).* arXiv:2303.17548.

Scherrer, N., Shi, C., Feder, A., and Blei, D. (2023). Evaluating the Moral Beliefs Encoded in LLMs (MoralChoice). *NeurIPS 2023.* arXiv:2307.14324.

Sorensen, T., Moore, J., Fisher, J., Gordon, M., et al. (2024). A Roadmap to Pluralistic Alignment. *International Conference on Machine Learning (ICML).* arXiv:2402.05070.

Srivastava, A., Rastogi, A., Rao, A., et al. (2022). Beyond the Imitation Game: Quantifying and Extrapolating the Capabilities of Language Models (BIG-bench). arXiv:2206.04615.

Wang, P., Li, L., Chen, L., Cai, Z., Zhu, D., Lin, B., Cao, Y., Kong, L., Liu, Q., Liu, T., and Sui, Z. (2023). Large Language Models are not Fair Evaluators. arXiv:2305.17926. (ACL 2024.)

Xu, C., Guan, S., Greene, D., and Kechadi, M.-T. (2024). Benchmark Data Contamination of Large Language Models: A Survey. arXiv:2406.04244.

Wataoka, K., Takahashi, T., and Ri, R. (2024). Self-Preference Bias in LLM-as-a-Judge. arXiv:2410.21819.

Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., and Stoica, I. (2023). Judging LLM-as-a-Judge with MT-Bench and Chatbot Arena. *NeurIPS 2023.* arXiv:2306.05685.

Westminster Assembly. (1647). *Westminster Confession of Faith; Larger Catechism; Shorter Catechism.* Repository-local primary-source transcriptions in `../WCF/`; no edition metadata asserted.